

Introduction to Audio Signal Processing with Librosa

- The video focuses on implementing the **amplitude envelope** feature from scratch using **librosa**, an audio processing library in Python ^s. It also covers plotting waveforms and the amplitude envelope itself ^s.
- `librosa` does not have a built-in amplitude envelope extractor, so one will be built manually ^s.

Loading and Playing Audio Files

- To begin, `librosa` and `librosa.display` (for plotting utilities) are imported ^s.
- Three 30-second WAV audio files are used for demonstration: an orchestral piece by Debussy, a jazz piece by Duke Ellington, and a rock song by Red Hot Chili Peppers ^s ^s.
- `IPython.display.Audio (ipd.Audio)` is used to play audio directly within a Jupyter Notebook environment ^s ^s.
- Audio files are loaded using `librosa.load()`, which returns the audio **signal** (a NumPy array) and its **sample rate** ^s.
- The default sample rate for `librosa.load()` is 22050 Hz, which is generally suitable for typical needs ^s.
- By default, `librosa.load()` converts audio to **mono** (single channel), which is often sufficient as significant information is usually not lost compared to stereo audio for many problems ^s.

Calculating the Amplitude Envelope

- The **amplitude envelope** for a specific frame is determined by taking the **maximum amplitude value** across all samples within that frame ^s.
- Two Python functions are demonstrated for calculating the amplitude envelope:

Simple Amplitude Envelope Function

- This function takes a `signal` and a `frame_size` ^s.
- It iterates through the signal in steps equal to the `frame_size`, effectively creating **non-overlapping frames** ^s ^s.
- For each frame, it calculates the maximum value of the signal slice `signal[i : i + frame_size]` ^s.
- The maximum values for each frame are collected into a list and then converted to a NumPy array ^s. A common `frame_size` is 1024 samples ^s.

Audio Signal Properties

- The total number of samples in a signal can be found using `signal.size` ^s .
- The **duration of one sample** (in seconds) is calculated as the inverse of the sample rate:
$$\text{sample_duration} = 1 / \text{sample_rate}$$
 ^s .
- The **total duration of the audio signal** (in seconds) is calculated by multiplying the sample duration by the total number of samples:
$$\text{signal_duration} = \text{sample_duration} \times \text{total_samples}$$
 ^s . For the examples, this confirms the 30-second length of each audio file ^s .

Visualizing Waveforms

- `matplotlib.pyplot` is imported as `plt` for plotting ^s .
- Waveforms are visualized using `librosa.display.waveplot()` (or `waveshow()`), which takes the signal as input ^s .
- Multiple waveforms can be stacked vertically using `plt.subplot(rows, columns, index)` ^s ^s .
- `plt.title()` sets the plot title, and `plt.ylim()` can set the y-axis range (e.g., -1 to 1 for normalized audio) ^s . Adding `alpha` can improve visual clarity ^s .

Waveform Characteristics by Genre

- **Debussy (Classical Music):** The waveform shows a fluid envelope with a large rise and fall of intensity, indicating high variability at a macro level due to acoustic instruments ^s ^s .
- **Red Hot Chili Peppers (Popular/Rock Music):** The waveform has a relatively stable overall envelope with distinct, regular spikes, typically caused by drum elements like kick and snare drums ^s ^s . This stability is common in popular music genres using electric instruments ^s .
- **Duke Ellington (Jazz Music):** Displays characteristics of both classical and rock music, with considerable intensity variability, though more at a micro level rather than large macro shifts ^s .

Music Genre	Waveform Characteristics	Amplitude Envelope Observations
Classical	Fluid envelope; high variability; large, macro-level changes in tension and energy.	Fluid; generally lower mean amplitude.
Rock/ Popular	Stable overall envelope; distinct, regular spikes (e.g., drum kit). Less variability.	Generally higher mean amplitude than classical or jazz; stable with clear, periodic spikes.
Jazz	Mix of both; considerable micro-level intensity variability.	Variable, but without the extreme macro-level changes of classical or the consistent high mean of popular music.

Amplitude Envelope with Overlapping Frames

- For **overlapping frames**, an additional parameter, `hop_length`, is introduced alongside `frame_size` .
- The `hop_length` specifies how many samples to shift to the right to start the next frame .
- The iteration step in the function is changed from `frame_size` to `hop_length` (`for i in range(0, signal.size, hop_length)`) .
- A typical `hop_length` might be 512 samples when `frame_size` is 1024 .

"Fancy" Amplitude Envelope Function

- A more concise "fancy" function achieves the same result using a single-line list comprehension within a NumPy array creation:

```
python
np.array([max(signal[i:i+frame_size]) for i in range(0, signal.size, hop_length)],
```

s s

- Both the simple and fancy functions produce identical results for the amplitude envelope calculation s s .

Visualizing the Amplitude Envelope

- The calculated amplitude envelope can be plotted alongside the waveform using `plt.plot()` .
- To align the envelope with the waveform's time axis, the frame indices of the amplitude envelope must be converted to time values s s .
- This conversion is done using `librosa.frames_to_time()`, which takes the frame indices (e.g., `np.arange(0, envelope.size)`) and the `hop_length` s .
- The plotted amplitude envelope (often in red) visually follows the overall intensity contour of the waveform s .
- **Key takeaways from visualization:** Rock/popular music often shows higher mean amplitude envelopes and distinct spikes from drums compared to classical or jazz music s .